

Ambient Knowledge

Ambient Knowledge is a Next.js demo of **AI-native message drafting**—where the AI is the core of the experience, but *stays out of your way*.

Instead of a chatbot, you get an “**ambient copilot**” **inside the composer**: as you type, it finds relevant internal context, explains *why* it was selected, and lets you insert grounded links.

Why it’s different

- **Invisible by design:** no prompt rituals, no separate chat—just better writing in-place.
- **Grounded:** suggestions always show sources (and you can insert /wiki/… links).
- **Trustable:** “Why this?” + scores/freshness, not magical autocomplete.
- **Resilient:** if LLM synthesis is unavailable, it continues in **retrieval-only** mode.

Quick start

```
npm install  
npm run dev
```

Optional LLM synthesis: - Copy `.env.example` → `.env.local` - Set `ANTHROPIC_API_KEY=...`

No key? Everything still works (retrieval-only + full source visibility).

60-second demo script

- 1) Open `http://localhost:3000`
- 2) Pick `#api-gateway`
- 3) Type:

Hey Sarah, can you approve the PR for the API gateway rate limiting rollout?
- 4) In the context card:
 - Open **Show details** to see **Serving tier**, **freshness**, and **Why this?** (top sources + scores)
 - Click **Insert link** → send → your message includes a clickable /wiki/… source
- 5) Click the source link to open the wiki page, then **Back to chat**.

Optional second scenario: - Go to `#platform-migration` - Type: “Hey Marcus, what is the status of the billing service migration?”

Notes

- This repo ships with mock knowledge data; no external services required for the core experience.
- When the LLM call fails (auth/rate limits), the UI stays responsive and continues surfacing context.